

1. MLP

```
In [ ]: # import library
import numpy as np

# Assign variables

# Weights
W1 = np.array([[1,-2],[3,4]])
W2 = np.array([[2,2],[2,-3]])

# Biases
b1 = np.array([[1],[0]])
b2 = np.array([[0],[-4]])

# MLP function
def mlpfunc(input, weight, bias):

    h = np.dot(weight, input)+bias    # Output for layer 1
    h[h<0] = 0    # ReLU

    return h

# Input
a0 = np.array([[1],[-1]])

a1 = mlpfunc(a0,W1,b1)
a2 = mlpfunc(a1,W2,b2)

# Output activation
print('Output activation:')
print(a2)

Output activation:
[[8]
 [4]]
```

2. HD5

```
In [ ]: import h5py
import numpy as np
import matplotlib.pyplot as plt

DEBUG = False
DATA_FNAME = 'WoojaeJeong_HW3_Prjb2.hd5'

if DEBUG:
    num_sequences = 3
    sequence_length = 4
else:
    num_sequences = 25
    sequence_length = 20

human_binary = np.array([
    [0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1],
    [1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0],
    [0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1],
    [0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0],
    [1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1],
    [1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0],
    [0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1],
    [1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1],
    [0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1],
    [0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1],
    [0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1],
    [1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1],
    [0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1],
    [0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1],
    [1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1],
    [0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0],
    [1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1],
    [1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0],
    [0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0],
    [1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1],
    [0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1],
    [1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1],
    [0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1],
    [0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1],
    [0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1],
    [0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1],
    [1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0],
    [0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0],
    [1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0],
    [0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0]
])

### do some error trapping:

assert human_binary.shape[0] == num_sequences, 'Error: the number of sequences was entered incorrectly'
assert human_binary.shape[1] == sequence_length, 'Error: the length of the sequences is incorrect'

# the with statement opens the file, does the business, and close it up for us...
with h5py.File(DATA_FNAME, 'w') as hf:
    hf.create_dataset('human_binary', data = human_binary)
    ## note you can write several data arrays into one hd5 file, just give each a different name.

#####
# Let's read it back from the file and then check to make sure it is as we wrote...
with h5py.File(DATA_FNAME, 'r') as hf:
    hb = hf['human_binary'][:]

### this will throw an error if they are not the same...
np.testing.assert_array_equal(human_binary, hb)

hb
```

```
Out [ ]: array([[0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1],
    [1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0],
    [0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1],
    [0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0],
    [1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1],
    [1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0],
    [0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1],
    [1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1],
    [0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1],
    [0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1],
    [0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1],
    [0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1],
    [1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1],
    [0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0],
    [1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1],
    [1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1],
    [0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0],
    [1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1],
    [0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1],
    [1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0],
    [0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1],
    [1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0],
    [0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0],
    [0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1],
    [1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1],
    [0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1],
    [1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0],
    [0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0]])
```

3. Logistic regression

(a) Logistic "2" detector

```
In [ ]: import h5py
import numpy as np
import matplotlib.pyplot as plt

# File name
Data_fName = 'mnist_traindata.hdf5'

# Desired number
DETECT_CLASS = 2

# Read train data
with h5py.File(Data_fName, 'r+') as df:
    xdata = df['xdata'][:]
    ydata = df['ydata'][:]

    # Find index for DETECT_CLASS (=2)
    target = (ydata.argmax(axis=1) == DETECT_CLASS)+0

# File name
Data_fName = 'mnist_testdata.hdf5'

# Desired number
DETECT_CLASS = 2

# Read train data
with h5py.File(Data_fName, 'r+') as df:
    xdata_test = df['xdata'][:]
    ydata_test = df['ydata'][:]

    # Find index for DETECT_CLASS (=2)
    target_test = (ydata_test.argmax(axis=1) == DETECT_CLASS)+0
```

```
In [ ]: # Logistic regression code

def sigmoid(data):
    return 1 / (1 + np.exp(-data))

# define probability
def p_x(data, weights, bias):
    px = sigmoid(data @ weights + bias)

    return px

def compute_cost(data, target, weights, bias, regcoef, regterm):

    px = p_x(data, weights, bias).squeeze()
    N = len(target)

    if regterm == 'l1':
        # l1 regularization
        regularization = regcoef*np.sum(np.abs(weights)).squeeze()

    if regterm == 'l2':
        # l2 regularization
        regularization = regcoef*np.sum(np.square(weights)).squeeze()

    if regterm == 'norm':
        regularization = 0

    cost = -(np.sum(target * np.log(px) + (1 - target) * np.log(1 - px)))/N + regularization

    return cost

def gradient_descent(data, target, weights, bias, regcoef, regterm, lrate, iterations, thresh):

    # Sample size
    N = len(target)
    prev_cost = float('inf')
    accuracy_data = []; loss_data = []

    for _ in range(iterations):

        # define p(x)
        px = sigmoid(np.dot(data, weights) + bias).squeeze()

        # Cost
        cost = compute_cost(data, target, weights, bias, regcoef, regterm)

        # Gradients for the weights
        if regterm == 'l1':
            weights_grad = np.dot(data.T, px - target)/N + regcoef*np.sign(weights).squeeze()

        if regterm == 'l2':
            weights_grad = np.dot(data.T, px - target)/N + regcoef*weights.squeeze()

        if regterm == 'norm':
            weights_grad = np.dot(data.T, px - target)/N

        bias_grad = np.sum(px-target)/N

        weights -= lrate * weights_grad
        bias -= lrate * bias_grad

        # Convergence threshold
        if abs(prev_cost - cost) < thresh: # if not, run until the maximum iteration
            break

        prev_cost = cost

        accuracy = np.sum((px >= 0.5)+0 == target)/len(data)
        accuracy_data.append(accuracy)
        loss_data.append(cost)
    return accuracy, cost, weights, bias, accuracy_data, loss_data
```

```
In [ ]: # Hyperparameters for GD without regularization

lrate = 0.1
iterations = 200
regcoef = 0.01
thresh = 1e-6
regterm = 'norm'

# Initialize weights and bias
weights = np.zeros((xdata.shape[1],1)).squeeze()
bias = 0

# Train the model
accuracy_norm, cost_norm, weights_norm, bias_norm, accuracy_data_norm, loss_data_norm = gradient_descent(xdata, target, weights, bias, regcoef, regterm, lrate, iterations, thresh)

# Test the model
accuracy_test_norm, cost_test_norm, weights_test_norm, bias_test_norm, accuracy_data_test_norm, loss_data_test_norm = gradient_descent(xdata_test, target_test, weights_norm, bias_norm,regcoef, regterm, lrate, iterations, thresh)
```

```
In [ ]: # Hyperparameters for GD wit l1 regularization

lrate = 0.1
iterations = 200
regcoef = 0.01
thresh = 1e-6
regterm = 'l1'

# Initialize weights and bias
weights = np.zeros((xdata.shape[1],1)).squeeze()
bias = 0

# Train the model
accuracy_l1, cost_l1, weights_l1, bias_l1, accuracy_data_l1, loss_data_l1 = gradient_descent(xdata, target, weights, bias, regcoef, regterm, lrate, iterations, thresh)

# Test the model
accuracy_test_l1, cost_test_l1, weights_test_l1, bias_test_l1, accuracy_data_test_l1, loss_data_test_l1 = gradient_descent(xdata_test, target_test, weights_l1, bias_l1,regcoef, regterm, lrate, iterations, thresh)
```

```
In [ ]: # Hyperparameters for GD wit l2 regularization

lrate = 0.1
iterations = 200
regcoef = 0.01
thresh = 1e-6
regterm = 'l2'

# Initialize weights and bias
weights = np.zeros((xdata.shape[1],1)).squeeze()
bias = 0

# Train the model
accuracy_l2, cost_l2, weights_l2, bias_l2, accuracy_data_l2, loss_data_l2 = gradient_descent(xdata, target, weights, bias, regcoef, regterm, lrate, iterations, thresh)

# Test the model
accuracy_test_l2, cost_test_l2, weights_test_l2, bias_test_l2, accuracy_data_test_l2, loss_data_test_l2 = gradient_descent(xdata_test, target_test, weights_l2, bias_l2,regcoef, regterm, lrate, iterations, thresh)
```

Answers


- i) I tried the learning rate for 0.01, 0.1, and 1. Lower the learning rate, it took longer to learn the features. I chose 0.1 which gave me relatively good results within a minute for 200 iterations.
- ii) I let the model to run for maximum iteration if it does not meet the convergence threshold.

```
In [ ]: it = np.arange(200)

fig = plt.figure(figsize = (12,4))

x = fig.add_subplot(121)
plt.plot(it, loss_data_norm)
plt.plot(it, loss_data_l1)
plt.plot(it, loss_data_l2)
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.legend(['Without', 'l1', 'l2'])

x = fig.add_subplot(122)
plt.plot(it, accuracy_data_norm)
plt.plot(it, accuracy_data_l1)
plt.plot(it, accuracy_data_l2)
plt.xlabel('Iterations')
plt.ylabel('Accuracy')
plt.legend(['Without', 'l1', 'l2'])
plt.show()
```



iii) I tried l1 and l2 regularization. With my set, implementing both l1, l2 regularization yielded higher loss and lower accuracy compared to the model without one. However, it prevents potential overfitting.

```
In [ ]: # Loss and accuracy plot for model with l2 regularization

it = np.arange(200)


fig = plt.figure(figsize = (12,4))

x = fig.add_subplot(121)
plt.plot(it, loss_data_l2)
plt.plot(it, loss_data_test_l2)
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.legend(['Training', 'Test'])

x = fig.add_subplot(122)
plt.plot(it, accuracy_data_l2)
plt.plot(it, accuracy_data_test_l2)
plt.xlabel('Iterations')
plt.ylabel('Accuracy')
plt.legend(['Training', 'Test'])
plt.show()

# Print out the final loss and finial accuracy for training and test set.
print('Training')
print('Accuracy =', accuracy_l2)
print('Loss =', cost_l2)

print('Test')
print('Accuracy =', accuracy_test_l2)
print('Loss =', cost_test_l2)
```



Training
Accuracy = 0.9645833333333333
Loss = 0.14801098072065189
Test
Accuracy = 0.973
Loss = 0.14235155961273654

```
In [ ]: # Create .hd5 files

DATA_FNAME = 'WoojaeJeong_HW3_Pr3a.hd5'
with h5py.File(DATA_FNAME, 'w') as hf:
    hf.create_dataset('w', data = np.asarray(weights_l2))
    hf.create_dataset('b', data = np.asarray(bias_l2))
```

(b) Softmax classification: gradient descent

```
In [ ]: import h5py
import numpy as np
import matplotlib.pyplot as plt

# load data
# File name
Data_fName = 'mnist_traindata.hdf5'

# Read train data
with h5py.File(Data_fName, 'r+') as df:
    xdata = df['xdata'][:,:]
    ydata = df['ydata'][:,:]

# File name
Data_fName = 'mnist_testdata.hdf5'

# Read train data
with h5py.File(Data_fName, 'r+') as df:
    xdata_test = df['xdata'][:,:]
    ydata_test = df['ydata'][:,:]

num_classes = 10
```

```
In [ ]: def softmax(data):
    z = np.exp(data - np.max(data))
    return z / z.sum(axis =1, keepdims=True)

def cost_softmax(data, target, weights, bias):
    N = len(target)
    px = softmax(np.dot(data, weights) + bias)
    cost = -np.sum(target * np.log(px + 10e-10)) / N

    return cost

def gradient_descent_softmax(xdata, ydata, weights, bias, lrate, iterations, thresh):
    # Sample size
    N = len(ydata)
    prev_cost = float('inf')

    accuracy_data = []; loss_data = []

    for _ in range(iterations):
        # define p(x)
        px = softmax(np.dot(xdata, weights) + bias)

        # Cost
        cost = cost_softmax(xdata, ydata, weights, bias)

        # Gradients for the weights
        weights_grad = np.dot(xdata.T, px - ydata)/N
        bias_grad = np.sum(px - ydata)/N

        weights -= lrate * weights_grad
        bias -= lrate * bias_grad

        # Convergence threshold
        if abs(prev_cost - cost) < thresh: # if not, run until the maximum iteration
            break

        prev_cost = cost

        accuracy = np.sum(np.argmax(px, axis=1) == np.argmax(ydata, axis=1))/N
        accuracy_data.append(accuracy)
        loss_data.append(cost)
    return accuracy, cost, weights, bias, accuracy_data, loss_data
```

```
In [ ]: # Hyperparameters

lrate = 0.1
iterations = 200
thresh = 1e-6

weights = np.zeros((xdata.shape[1],num_classes))
bias = np.zeros((1,num_classes))

# Train the model
accuracy, cost, weights_train, bias_train, accuracy_data, loss_data = gradient_descent_softmax(xdata, ydata, weights, bias, lrate, iterations, thresh)

# Test the model
accuracy_test, cost_test, weights_test, bias_test, accuracy_data_test, loss_data_test = gradient_descent_softmax(xdata_test, ydata_test, weights, bias, lrate, iterations, thresh)
```

(ii) I chose the learning rate = 0.1

```
In [ ]: it = np.arange(200)


fig = plt.figure(figsize = (12,4))

x = fig.add_subplot(121)
plt.plot(it, loss_data)
plt.plot(it, loss_data_test)
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.legend(['Training','Test'])

x = fig.add_subplot(122)
plt.plot(it, accuracy_data)
plt.plot(it, accuracy_data_test)
plt.xlabel('Iterations')
plt.ylabel('Accuracy')
plt.legend(['Training','Test'])
plt.show()

# Print out the final loss and finial accuracy for training and test set.
print('Training')
print('Accuracy =', accuracy)
print('Loss =', cost)

print('Test')
print('Accuracy =', accuracy_test)
print('Loss =', cost_test)
```



Training
Accuracy = 0.8759333333333333
Loss = 0.4925534953935439
Test
Accuracy = 0.9
Loss = 0.3797783969872003

(c) Softmax classification: stochastic gradient descent

```
In [ ]: import h5py
import numpy as np
import matplotlib.pyplot as plt

# load data
# File name
Data_fName = 'mnist_traindata.hdf5'

# Read train data
with h5py.File(Data_fName, 'r+') as df:
    xdata = df['xdata'][:, :]
    ydata = df['ydata'][:, :]

# File name
Data_fName = 'mnist_testdata.hdf5'

# Read train data
with h5py.File(Data_fName, 'r+') as df:
    xdata_test = df['xdata'][:, :]
    ydata_test = df['ydata'][:, :]

num_classes = 10
```

```
In [ ]: def softmax(data):
    z = np.exp(data - np.max(data))
    return z / z.sum(axis =1, keepdims=True)

def cost_softmax(data, target, weights, bias):
    N = len(target)
    px = softmax(np.dot(data, weights) + bias)
    cost = -np.sum(target * np.log(px + 10e-10)) / N

    return cost

def SGD_softmax(xdata, ydata, weights, bias, lrate, iterations, batch_size):

    # Sample size
    N = len(ydata)
    prev_cost = float('inf')

    accuracy_batch = []; loss_batch = []
    batch_plot = 5000

    for _ in range(iterations):
        for idx in range(0, N, batch_size):

            batch_end_idx = min(idx + batch_size, N)

            x_batch, y_batch = xdata[idx:batch_end_idx], ydata[idx:batch_end_idx]

            # define p(x)
            px = softmax(np.dot(x_batch, weights) + bias)

            # Gradients for the weights
            weights_grad = -np.dot(x_batch.T, y_batch - px)/batch_size
            bias_grad = -np.sum(y_batch - px, axis = 0)/batch_size

            weights -= lrate * weights_grad
            bias -= lrate * bias_grad

            if idx % batch_plot == 0:
                # Cost
                cost = cost_softmax(xdata, ydata, weights, bias)

                # Accuracy
                px = softmax(np.dot(xdata, weights) + bias)
                accuracy = np.mean(np.argmax(px, axis=1) == np.argmax(ydata, axis=1))

                accuracy_batch.append(accuracy)
                loss_batch.append(cost)

    return accuracy, cost, weights, bias, accuracy_batch, loss_batch
```

```
In [ ]: # Hyperparameters

iterations = 1
batch_size = 1

# Initial weights and bias
weights = np.zeros((xdata.shape[1],num_classes))
bias = np.zeros((1,num_classes))

# Train the model with 1e-2 learning rate
accuracy, cost, weights_train, bias_train, accuracy_batch, loss_batch = SGD_softmax(xdata, ydata, weights, bias, 1e-2, iterations, batch_size)

# Test the model with 1e-2 learning rate
iterations = 6
accuracy_test, cost_test, weights_train_test, bias_train_test, accuracy_batch_test, loss_batch_test = SGD_softmax(xdata_test, ydata_test, weights_train, bias_train, 1e-2, iterations, batch_size)

# Initial weights and bias
weights = np.zeros((xdata.shape[1],num_classes))
bias = np.zeros((1,num_classes))

# Train the model with 1e-4 learning rate
iterations = 1
accuracy2, cost2, weights_train2, bias_train2, accuracy_batch2, loss_batch2 = SGD_softmax(xdata, ydata, weights, bias, 1e-4, iterations, batch_size)

# Test the model with 1e-4 learning rate
iterations = 6
accuracy_test2, cost_test2, weights_train_test2, bias_train_test2, accuracy_batch_test2, loss_batch_test2 = SGD_softmax(xdata_test, ydata_test, weights_train2, bias_train2, 1e-4, iterations, batch_size)
```

```
In [ ]: batch_number = np.arange(0, len(ydata), 5000/batch_size)


fig = plt.figure(figsize = (12,4))

x = fig.add_subplot(121)
plt.plot(batch_number,accuracy_batch)
```

```
plt.plot(batch_number,accuracy_batch2)
plt.plot(batch_number,loss_batch)
plt.plot(batch_number,loss_batch2)
plt.xlabel('Batch number')
plt.legend(['Accuracy (Learning rate= 1e-2)', 'Accuracy (Learning rate= 1e-4)', 'Loss (Learning rate= 1e-2)', 'Loss (Learning rate= 1e-4)'])
plt.title('Training')

batch_number = np.arange(0, len(ydata_test)*6, 5000/batch_size)

x = fig.add_subplot(122)
plt.plot(batch_number,accuracy_batch_test)
plt.plot(batch_number,accuracy_batch_test2)
plt.plot(batch_number,loss_batch_test)
plt.plot(batch_number,loss_batch_test2)
plt.xlabel('Batch number')
plt.legend(['Accuracy (Learning rate= 1e-2)', 'Accuracy (Learning rate= 1e-4)', 'Loss (Learning rate= 1e-2)', 'Loss (Learning rate= 1e-4)'])
plt.title('Test')
plt.show()
```



In []: `print(accuracy)`

0.8976833333333334

(i) With learning rate = 0.01, it took less than 10,000 iterations to acheive the comparable performance with batch gradient descent (~0.88). With learning rate = 0.0001, it took more than 150,000 iterations to acheive the comparable performance. With higher learning rate, the model learned faster.

(ii) To acheive 88% accuracy in the training set, the batch gradient-descent model took 200 iterations. So, the computational complexity is $60,000 \times 200 = 12,000,000$. To acheive the comparable performance, the SGD took less than 10,000 iterations with a single-batch. So, the computational complexity is $10,000 \times 1 = 10,000$. The computational complexity is less in SGD.

```
# Hyperparameters for mini-batch size of 100
iterations = 10
batch_size = 100

# Initial weights and bias
weights = np.zeros((xdata.shape[1],num_classes))
bias = np.zeros((1,num_classes))

# Train the model with 1e-2 learning rate
accuracy, cost, weights_train, bias_train, accuracy_batch, loss_batch = SGD_softmax(xdata, ydata, weights, bias, 1e-2, iterations, batch_size)

# Test the model with 1e-2 learning rate
iterations = 60
accuracy_test, cost_test, weights_train_test, bias_train_test, accuracy_batch_test, loss_batch_test = SGD_softmax(xdata_test, ydata_test, weights_train, bias_train, 1e-2, iterations, batch_size)

# Initial weights and bias
weights = np.zeros((xdata.shape[1],num_classes))
bias = np.zeros((1,num_classes))

# Train the model with 1e-4 learning rate
iterations = 10
accuracy2, cost2, weights_train2, bias_train2, accuracy_batch2, loss_batch2 = SGD_softmax(xdata, ydata, weights, bias, 1e-4, iterations, batch_size)

# Test the model with 1e-4 learning rate
iterations = 60
accuracy_test2, cost_test2, weights_train_test2, bias_train_test2, accuracy_batch_test2, loss_batch_test2 = SGD_softmax(xdata_test, ydata_test, weights_train2, bias_train2, 1e-4, iterations, batch_size)
```


```
In [ ]: batch_number = np.arange(0, len(ydata)*10, 5000)

fig = plt.figure(figsize = (12,4))

x = fig.add_subplot(121)
plt.plot(batch_number,accuracy_batch)
plt.plot(batch_number,accuracy_batch2)
plt.plot(batch_number,loss_batch)
plt.plot(batch_number,loss_batch2)
plt.xlabel('Batch number')
plt.legend(['Accuracy (Learning rate= 1e-2)', 'Accuracy (Learning rate= 1e-4)', 'Loss (Learning rate= 1e-2)', 'Loss (Learning rate= 1e-4)'])
plt.title('Training')

batch_number = np.arange(0, len(ydata_test)*60, 5000)

x = fig.add_subplot(122)
plt.plot(batch_number,accuracy_batch_test)
plt.plot(batch_number,accuracy_batch_test2)
plt.plot(batch_number,loss_batch_test)
plt.plot(batch_number,loss_batch_test2)
plt.xlabel('Batch number')
plt.legend(['Accuracy (Learning rate= 1e-2)', 'Accuracy (Learning rate= 1e-4)', 'Loss (Learning rate= 1e-2)', 'Loss (Learning rate= 1e-4)'])
plt.title('Test')
plt.show()
```



(iii) With learning rate = 0.01, it took around 150,000 iterations to acheive the comparable performance with batch gradient descent (~0.88). With learning rate = 0.0001, it did not acheive the comparable performance up to 6,000,000 iterations. With higher learning rate, the model learned faster. To acheive 88% accuracy in the training set, the batch gradient-descent model took 200 iterations. So, the computational complexity is $60,000 \times 200 = 12,000,000$. To acheive the comparable performance, the SGD took around 150,000 iterations with a mini-batch with size = 100. So, the computational complexity is $150,000 \times 100 = 15,000,000$. The computational complexity is less in SGD.

(iv) Overall, to acheive ~88% performance, the computational complexity: 1) 100 sample mini-batch = 1,500,000, 2) the single-point mini-batch = 10,000, and 3) batch gradient descent = 12,000,000. From the experiment, the time to train the model to ~88% was shortest with the single-point mini-batch.

```
In [ ]: # Create .hd5 files
w_train = weights_train.T
b_train = bias_train.T.squeeze()

DATA_FNAME = 'WoojaeJeong_HW3_Pr33c.hd5'
with h5py.File(DATA_FNAME, 'w') as hf:
    hf.create_dataset('w', data = np.asarray(w_train))
    hf.create_dataset('b', data = np.asarray(b_train))
```